

6 Activation Records

( 6. Activation Records , p. 125)


- Tiger allows the use of **recursive** function calls and the use of **local variables** within a function.
 - **Example:** compute the greatest common divisor of m and n (we assume the presence of a function mod):

Recursive Tiger function gcd

```
1  function gcd (m : int, n : int) : int =  
2      let var x := mod (m, n)  
3      in  
4          if n = 0 then m  
5              else gcd (n, x)  
6      end
```

- We expect each invocation of gcd to be **independent** of all others: each invocation has its own copy of variables m , n , and x .

- Since many such invocations may be active concurrently, we can *not* use some global memory location to hold these variables.
 - **Example: call graph** for invocation `gcd (21,12)`:

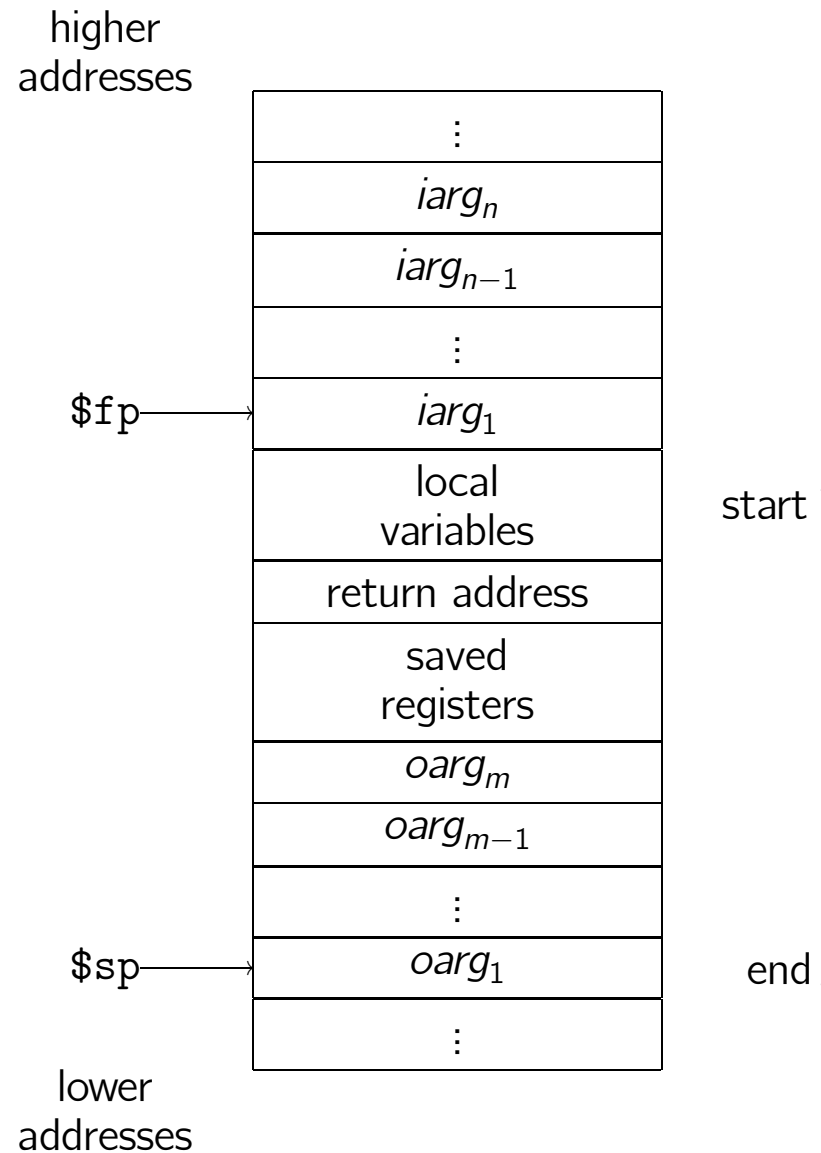


- Each **invocation needs to be assigned its own local variable storage** that is released in a **LIFO** manner when the function returns.
- In Tiger (as is most programming languages), this local variable storage is placed in the CPU **stack**. The depth of the call graph is thus bounded by the stack size (typically, the stack may grow until it fills the entire available heap memory).

6.1 Stack Frames

- Consider an invocation of function `gcd`:
 - Whenever `gcd` is called, we need to grow the stack by 3 entries (*push* arguments `m`, `n`, and local variable `x`, although `x`'s value is yet unknown – what should we *push* there?).
 - Similarly, when `gcd` returns we need to dispose the 3 topmost stack entries in one batch ($3 \times \text{pop}$).
- So, instead of the “classical” *push/pop* stack model, we rather need a technique to grow/shrink the stack in batches:
 - We use the stack like a dynamic array and adjust the **stack pointer** (MIPS: CPU register `$sp`) for each invocation only once.
 - Stack entries *beyond* `$sp` are garbage, entries *before* `$sp` are considered allocated.
 - On function invocation, we decrease `$sp` so that function arguments and local variables will fit in the newly allocated area (the **stack frame/activation record**).

- A typical stack frame layout:²²



²²CPU vendors often prescribe a standard frame layout to ensure interoperability of different programming languages.

- In the following, we now turn to and discuss each of the depicted frame slots.

6.1.1 The Frame Pointer

- When a function $g()$ calls out to a function $f(a_1, \dots, a_n)$, we refer to g as the **caller** and f as the **callee**.
- What happens during a call?
 - ① On entry to f , $\$sp$ points to the first **incoming argument** a_1 .
 - ② f allocates a frame on the stack by subtracting the needed **frame size** s from $\$sp$.
 - ③ The current value of the **frame pointer** (MIPS: $\$fp$) is stored in the frame (slot *saved registers*). The old value of $\$sp$ is remembered in register $\$fp$.
- What happens when function f exits?
 - ① Copy $\$fp$ to $\$sp$ (discards the frame).
 - ② Restore the old $\$fp$ value from the just discarded frame.

6.1.2 Registers

- Modern CPUs, especially RISC CPUs, come with a large number (MIPS: 32) of CPU internal **registers**.
 - Register read/write access is much faster than any memory read/write using *load/store* instructions (also faster than CPU cache memory).
 - An optimizing compiler will thus try to keep local variables, intermediate evaluation results, etc. in CPU registers.
 - Given that many function invocations can be active in parallel (depth of the call graph), registers are a scarce resource.
 - **Example:**
 - *g* uses register *r* and calls out to *f*.
 - *f* needs to store *r* in its stack frame before it may write to *r* and later restores *r* from the frame before *f* returns (*r* is **callee-saved**)

OR

g stores *r* in its frame, then calls *f* which clobbers *r* at will until return to *g*, then *g* restores *r* from its own frame (*r* is **caller-saved**).

- CPU vendors usually introduce conventions about CPU registers usage.

Example: MIPS:

MIPS Register	Usage
\$zero	constant 0
\$at	reserved for assembler
\$v0	expression evaluation and function results
\$v1	expression evaluation and function results
\$a0	argument 1
:	:
\$a3	argument 4
\$t0	temporary (caller-save)
:	:
\$t7	temporary (caller-save)
\$s0	saved temporary (callee-save)
:	:
\$s7	saved temporary (callee-save)
\$t8	temporary (caller-save)
\$t9	temporary (caller-save)
\$k0	reserved for OS kernel
\$k1	reserved for OS kernel
\$gp	pointer to global area
\$sp	stack pointer
\$fp	frame pointer
\$ra	return address (used by function call)

- Sometimes, however, caller and or callee can do without register saving.

Examples:


- If g (caller) knows it will not **need** variable x after calling f , hold x in a caller-save register and *do not save r* .
 - If g knows it will **need** variable x after several function calls have been made, hold x in a caller-save register r and save/restore r just once.
- The “wise” assignment of variables to caller/callee-save registers is an important compiler optimization (backed up by **data-flow analysis** techniques).

6.1.3 Parameter Passing

- In the frame layout shown before, a function can address the different entries in its frame **relative** to the **frame pointer** $\$fp$ and the **stack pointer** $\$sp$:
 - The function may access its **incoming arguments** $iarg_1, \dots, iarg_k$ (if the function is of arity k) at or just above $\$fp$.
 - Likewise, if the function calls out, it places **outgoing arguments** in frame slots $oarg_1 \dots oarg_m$ at or just above $\$sp$.²³
- Stack-based parameter passing causes significant memory traffic.
- Modern CPU architectures thus reserve a fixed set of k , say, **registers for parameter passing** (MIPS: $k = 4$, $\$a0 \dots \$a3$). If a function takes $n > k$ parameters, only parameters $k + 1 \dots n$ are passed via the stack.


²³Note that $oarg_i$ will be $iarg_i$ for the called function.

- By definition, argument passing registers need to be operated in caller-save mode:

-  Suppose $g(a_1, \dots, a_n)$ receives a_1 in $\$a0$. By convention, if g calls $f(x)$, it needs to pass x in $\$a0$, so g must save $\$a0$ (in its stack frame).


So, how is any memory traffic saved?

- ① Some functions are **leaf functions**, i.e., they are leaves in the program's call graph:



Leaf functions *never* need to caller-save arguments. More importantly, leaf functions often *need no stack frame at all*.

- ② Some compilers use **interprocedural register allocation**, e.g., g might receive a_1 in $\$a0$ but pass x to $f(x)$ in $\$a3$ (and the code of f is prepared accordingly).
- ③ g might not be a leaf procedure but may be done with its use of argument register $\$a0$ before it calls out to f .

-  Parameter passing via argument registers introduces problems, when the callee needs to access the **address** of an argument.

Examples:

- The varargs feature of C, e.g., used in `printf`.
- Taking the address of an argument a_i , e.g. in C:

C code fragment

```
1      int *f (int x)
2      {
3          return &x;
4      }
```

6.1.4 Return Address

- When the callee f has completed its work, its last job is to ensure that **execution is resumed** just after the associated $f(\dots)$ call located in f 's caller g .
 - Resuming is implemented by loading the CPU's instruction pointer with the **return address** ra :

```

1      function g (...) =
2          (s1; ...; f(...); sk; ...; sn)
                                ↑
                                ra

```

- Modern CPUs save *ra* in a register rather than on the stack.
 - Non-leaf functions may then save *ra* in their stack frame only if actually needed.
(MIPS: *ra* is automatically saved in register \$ra by instruction *jal* (*jump and link*) , the callee resumes by executing the *jump* instruction

j \$ra .


MIPS function calling and return

```
1      g:  :  
2          jal f  
3          :  
4  
5      f:  :  
6          j $ra
```

6.1.5 Frame-Resident Variables

- Function calling conventions try to ensure to generate as few memory traffic as possible.
- There are some circumstances where a function cannot avoid writing variable values, temporary expression evaluation results, etc., to the stack frame, however:
 - ① The value is **too big** to fit into a single CPU register.
 - ② The register currently holding the value is needed for a different purpose (**caller/callee saving**, see previous subsections).
 - ③ There are too many variables and temporary results that will not all fit into CPU registers (some of these thus need to be **spilled** into the stack frame).
 - ④ The variable holding the value is accessed by a function **nested inside** the current one (\rightarrow *static links*, see below). The variable **escapes**.

- One strategy to compile a program would be to **assign each variable a location** (either a register or a slot in the stack frame) **as soon as it is declared**.

-  This is *not* feasible, however, since the compiler will not yet know which registers will be available, how many local variables and temporary expression results will be needed, if a variable will be accessed from inside a nested function etc.


Play safe: assign (provisional) locations for *all* parameters, local variables, and temporary results *in the stack frame*. An optimization phase may later “move” in-frame locations to registers.

This optimization phase is commonly referred to as **register allocation**.

6.1.6 Static Links

- Tiger (like Pascal) allows for the declaration of functions **nested inside** other functions. Consequently, a source program imposes a **static hierarchical structure** on the function declarations contained in it.

Example: for the Tiger program given on the next slide, the static hierarchical structure is:




- The **scoping rules** of Tiger permit that a function in this hierarchy may **access the parameters and variables of any of its ancestors**.

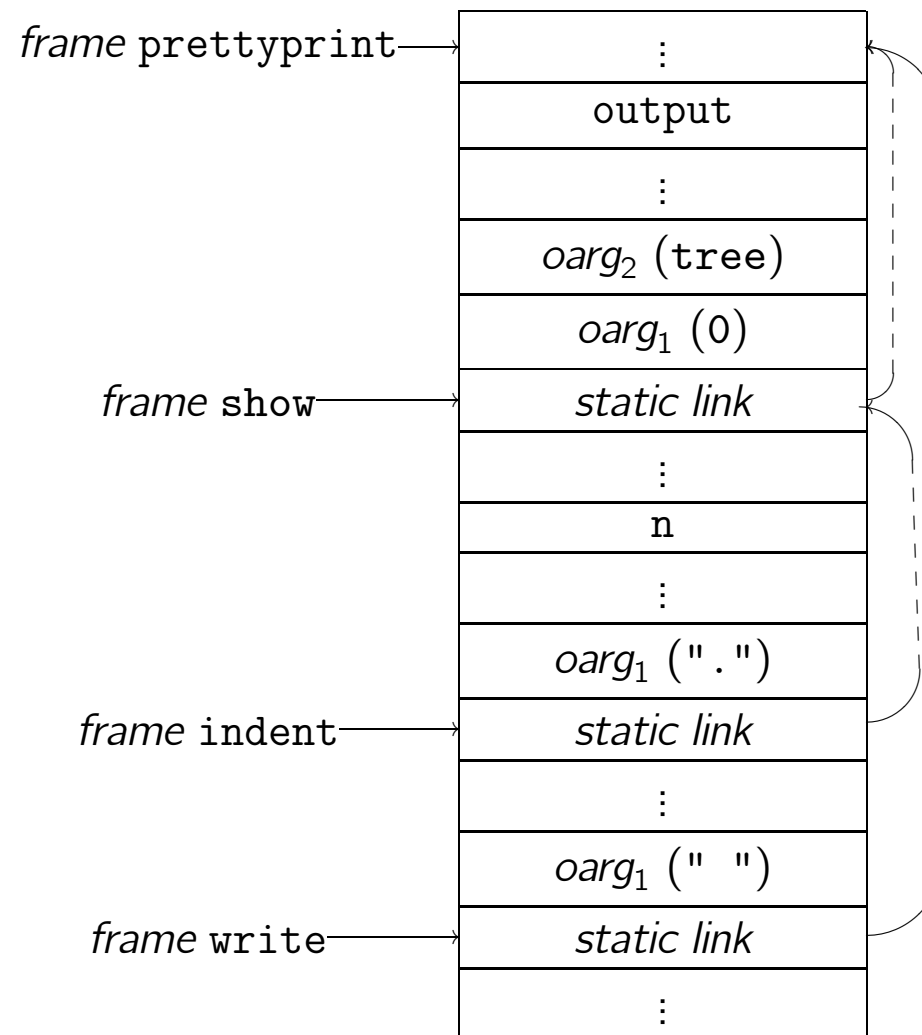
Tiger program with nested functions


```
1  type tree = { key: string, left: tree, right: tree }
2
3  function prettyprint (tree: tree) : string =
4      let
5          var output := ""
6
7          function write (s: string) =
8              output := concat (output, s)
9
10         function show (n: int, t: tree) =
11             let function indent (s: string) =
12                 (for i := 1 to n do
13                     write (" ");
14                     output := concat (output, s); write ("\n"))
15             in if t = nil then indent (".")
16                else (indent (t.key);
17                     show (n + 1, t.left);
18                     show (n + 1, t.right ))
19         end
20
21     in show (0, tree); output
22 end
```

- This program contains several references to non-local variables of functions higher up in the static hierarchy:
 - ① line 8: write references output (output **escapes** from prettyprint).
 - ② line 12: indent references n (n **escapes** from show).
 - ③ line 14: indent references output.
- To implement references to escaped variables, we make sure that each called function f finds the **frame pointer of its parent in the static hierarchy** in a predefined slot in its own frame. This is called the **static link**.
 - If f needs to access a non-local variable of any ancestor g higher up in the hierarchy, it simply follows the chain of static link pointers to find g 's frame pointer.
 - *length of chain = difference of static nesting depths of g and f :*



- State of the CPU stack after calling sequence `prettyprint-show-indent-write`:



-  The static link of `write`'s activation record points to the frame of its static parent `prettyprint`, *not* the frame of its “dynamic parent” (= caller) `indent`.

Examples:

- indent uses show's escaped parameter n:

$$depth_{\text{indent}} - depth_{\text{show}} = 2 - 1 = 1 .$$

The compiler thus needs to insert code to follow the static link chain *once* to obtain show's frame pointer which is then used to access n.

- write uses prettyprint's escaped variable output:

$$depth_{\text{write}} - depth_{\text{prettyprint}} = 1 - 0 = 1 .$$

Insert code to follow the static link chain *once* to obtain prettyprint's frame pointer which is then used to access output.

- indent uses prettyprint's escaped variable output:

$$depth_{\text{indent}} - depth_{\text{prettyprint}} = 2 - 0 = 2 .$$

Insert code to follow the static link chain *twice* to obtain prettyprint's frame pointer which is then used to access output.

6.2 Frames in the Tiger Compiler

- The actual frame layout our compiler needs to use depends on a number of issues, among these notably
 - ① the target CPU type,
 - ② and programming language interoperability conventions (suppose we want to interface Tiger programs with a C library).
- To prevent such details from invading the semantic analysis phase of the compiler, we instead rely on an **abstract frame interface** (specified in C header file `frame.h`).
 - Each target-specific backend then only needs to implement this well-defined frame interface (e.g., we can have implementations like `mipsframe.c`, `pentiumframe.c`, `sparcframe.c`).

- The C data structure representing a frame remains abstract (i.e., the actual C `struct F_frame_ { ... }` is only found in the target-specific `mipsframe.c`, ...):

_____ frame.h **(to be continued)** _____

```
1 typedef struct F_frame_ *F_frame;
```

- The following function `F_newFrame` creates a new frame:

_____ frame.h **(continued)** _____

```
2 F_frame F_newFrame (Temp_label name, U_boolList formals);
```

- `name` denotes the function name²⁴,
- `formals` is a list of k Booleans (if the function is k -ary). If the n -th entry in this list is TRUE, the n -th parameter of the function escapes, otherwise it does not.

Example: the frame for function `show` (`n: int`, `t: tree`) with `n` escaping from `show`, could be created via

```
f = F_newFrame (Temp_namelabel ("show"),
                U_BoolList (TRUE /* n */, U_BoolList (FALSE /* t */, NULL)));
```

²⁴For now, think of `Temp_label` as of other symbols used in the compiler.

- Whenever a local variable needs to be allocated in the frame:

```
frame.h (continued)  
5 F_access F_allocLocal (F_frame f, bool escape);
```

- f denotes the frame we need to allocate the variable in,
- $escape$ indicates if this variable needs to be allocated in frame or may live in a temporary register ($escape = false$).
- $F_allocLocal$ returns the **access** information needed to reference the new variable later on. An access is either
 - ① $InFrame(o)$, where o is an offset relative to the frame pointer,
 - ② $InReg(t)$, with t denoting a **temporary location**²⁵. In this compilation phase we simply assume an infinite pool of temporaries.

Example: on the MIPS, successive $F_allocLocal(f, TRUE)$ calls yield accesses of the form

$InFrame(-4), InFrame(-8), InFrame(-12), \dots$

because the **word size** W for the MIPS architecture is 4.

²⁵Register allocation will later replace t by a specific CPU register.

- Likewise, the access information for all formal parameters of a function is returned by

```
_____ frame.h (continued) _____  
6  F_accessList F_formals (F_frame f);
```

Example: for the frame `f` of function `show (n: int, t: tree)` with `n` escaping, `F_formals (f)` might return the list

`[InFrame (0), InReg (t_2)]`

or, if we conservatively assume that all formals escape:

`[InFrame (0), InFrame (4)]`

- Lastly, the implementation of the frame module (e.g., `mipsframe.c`) needs to encapsulate details of the **calling conventions** used on the target machine.

Example:

- MIPS conventions demand that caller passes the first four parameters in `$a0 ... $a3`.
- In a simple compiler we might however choose to use frame-resident variables only (i.e., all accesses are of the form `InFrame(o)`).
- The frame module provides a *sequence of instructions*, the **view shift**, ensuring that, on function entry, parameters are moved into the callee's frame where they are expected by the callee's body.

Example: MIPS view shift for a two-argument function (parameters passed by caller in `$a0`, `$a1`) with two frame-resident local variables (frame size s , $M[a]$: memory access at address a):

MIPS view shift (pseudo code)	
1	<code>\$fp ← \$sp</code>
2	<code>\$sp ← \$sp - s</code>
3	<code>M[fp-12] ← \$a0</code>
4	<code>M[fp-16] ← \$a1</code>